Probabilistic time series forecasting with Facebook prophet

Sam Leach
github / twitter: samuelleach

PyData London 2017 lightning talk
https://github.com/samuelleach/python-notebooks

Old problem

Something happened yesterday.

Will it happen tomorrow?

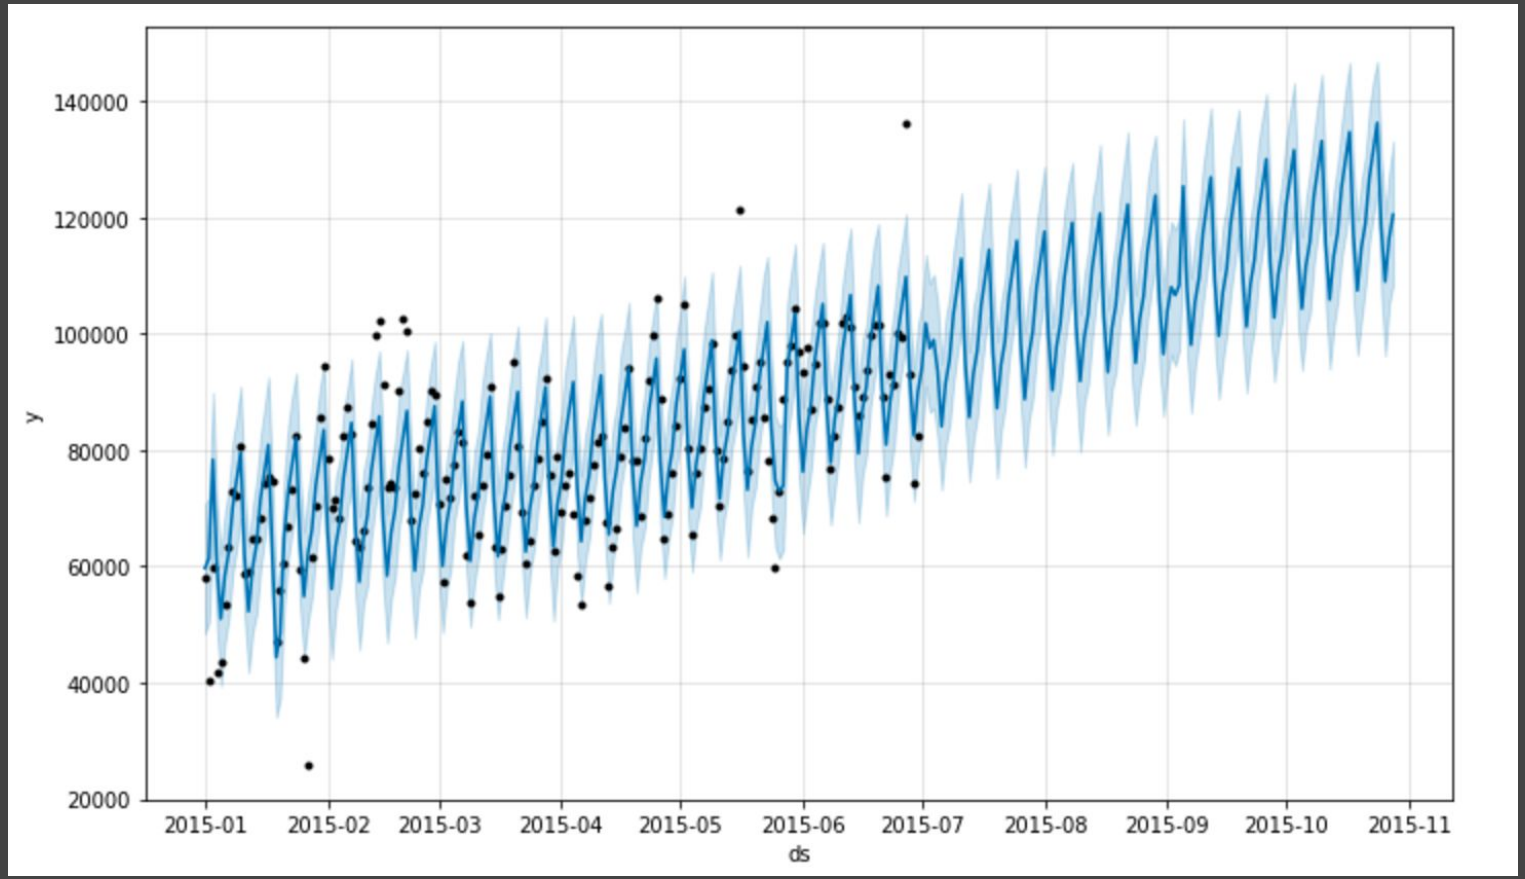
140000
120000
100000
80000
60000
40000
20000
y
2015-01
2015-02
2015-03
2015-04
2015-05
2015-06
2015-07
2015-08
2015-09
2015-10
2015-11
ds

PROPHET

Docs GitHub

Forecasting at scale.

Prophet is a forecasting procedure implemented in R and Python. It is fast and provides completely automated forecasts that can be tuned by hand by data scientists and analysts.

INSTALL PROPHET | GET STARTED IN R | GET STARTED IN PYTHON | READ THE PAPER

Prophet is a procedure for forecasting time series data. It is based on an additive model where non-linear trends are fit with yearly and weekly seasonality, plus holidays. It works best with daily periodicity data with at least one year of historical data. Prophet is robust to missing data, shifts in the trend, and large outliers.

Prophet is open source software released by Facebook's Core Data Science team. It is available for download on CRAN and PyPI.

Accurate and fast.

Prophet is used in many applications across Facebook for producing reliable forecasts for planning and goal setting. We've found it to perform better than any other approach in the majority of cases. We fit models in Stan so that you get forecasts in just a few seconds.

Fully automatic.

Get a reasonable forecast on messy data with no manual effort. Prophet is robust to outliers, missing data, and dramatic changes in your time series.

Tunable forecasts.

The Prophet procedure includes many possibilities for users to tweak and adjust forecasts. You can use human-interpretable parameters to improve your forecast by adding your domain knowledge.

Available in R or Python.

We've implemented the Prophet procedure in R and Python, but they share the same underlying Stan code for fitting. Use whatever language you're comfortable with to get forecasts.

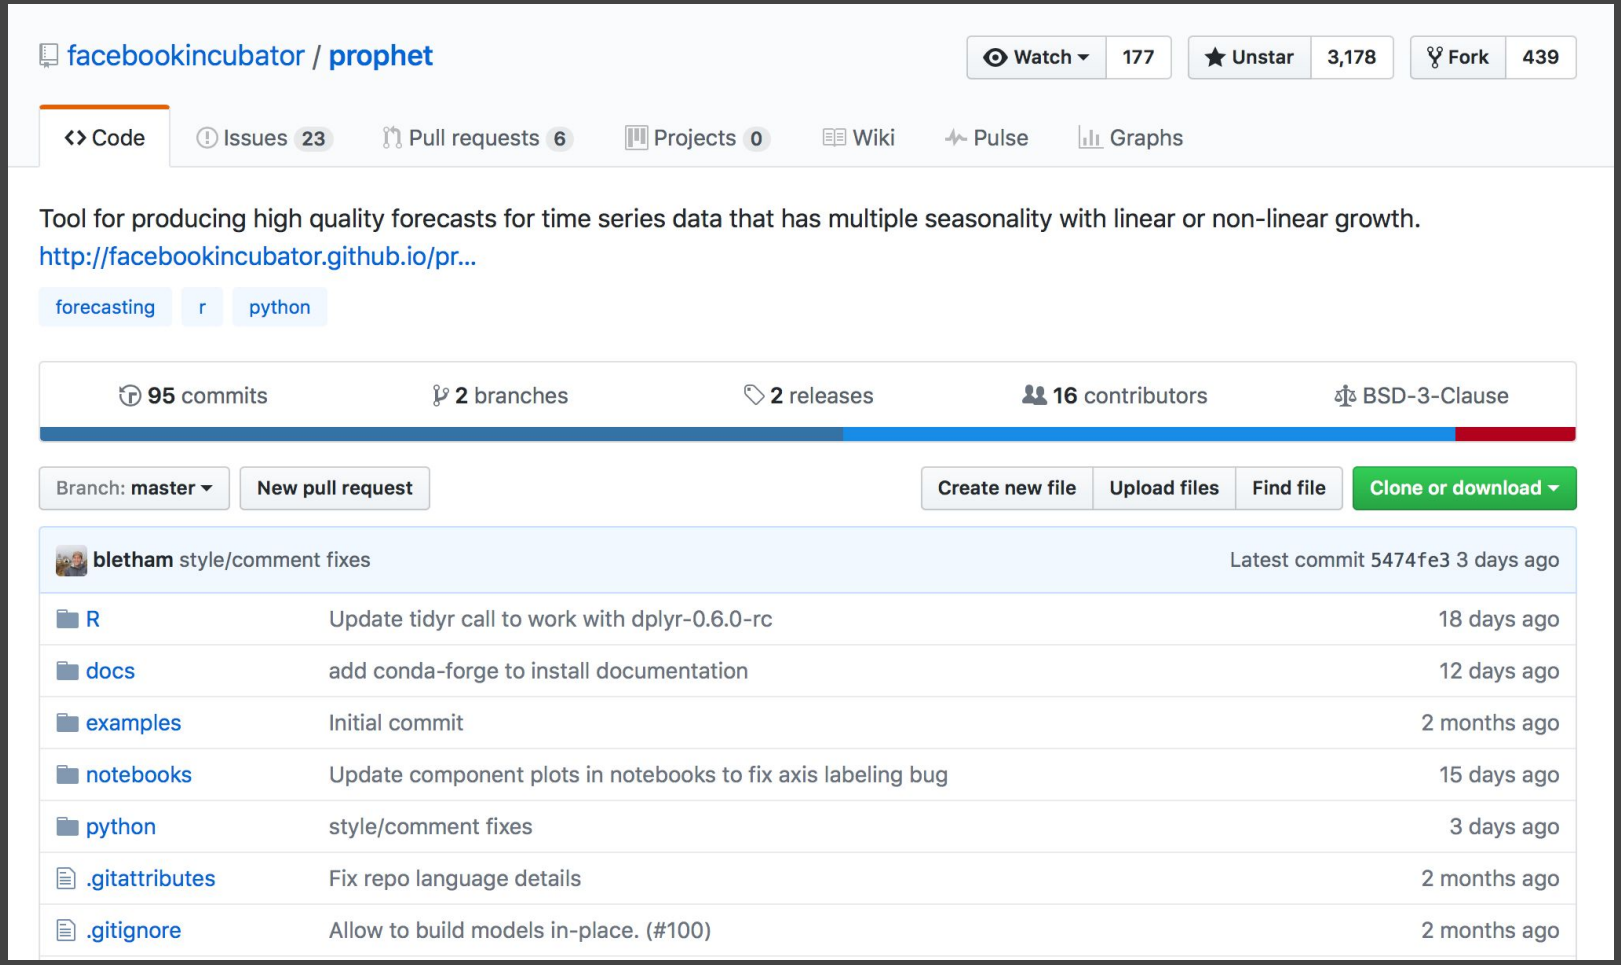
facebookincubator / prophet
Watch 177
Unstar 3,178
Fork 439
Code
Issues 23
Pull requests 6
Projects 0
Wiki
Pulse
Graphs
Tool for producing high quality forecasts for time series data that has multiple seasonality with linear or non-linear growth.
http://facebookincubator.github.io/pr...
forecasting
r
python
95 commits
2 branches
2 releases
16 contributors
BSD-3-Clause
Branch: master
New pull request
Create new file
Upload files
Find file
Clone or download
bletham style/comment fixes
Latest commit 5474fe3 3 days ago
R
Update tidyr call to work with dplyr-0.6.0-rc
18 days ago
docs
add conda-forge to install documentation
12 days ago
examples
Initial commit
2 months ago
notebooks
Update component plots in notebooks to fix axis labeling bug
15 days ago
python
style/comment fixes
3 days ago
.gitattributes
Fix repo language details
2 months ago
.gitignore
Allow to build models in-place. (#100)
2 months ago

What is Prophet?

- Time series forecasting library by Sean Taylor and Benjamin Letham at Facebook.
- Models a single time series with a daily timestamp.
- Scikit learn like fit / predict API.
- Uses stan under the hood to sample uncertainties of some model component parameters.

Time series model components

- Linear or non-linear growth with changepoint detection and specification.
- Weekly and yearly periodic / Fourier components.
- Specification of (multiple groups of) holiday events.

Inputs

- Pandas dataframe with a daily timestamp and time series to be modeled.
- Model parameters: changepoints, holidays.

Outputs

- Dataframe with forecasts and uncertainties for some set of future timestamps.
- Model components and uncertainties.

Demo

Run prophet on 14 million Uber rides in New York, January - June 2015

Notebook: https://github.com/samuelleach/python-notebooks
Data: https://github.com/fivethirtyeight/uber-tlc-foil-response

In [14]:
```
# Second model including uncertainties in the weekly component
# and holidays, sampled via MCMC.

model = Prophet(mcmc_samples=500, holidays=federalholidays, yearly_seasonality=False)
```

In [15]:
```
model.fit(df_agg)

future = model.make_future_dataframe(periods=30)

forecast = model.predict(future)
model.plot(forecast)
```

Out[15]:

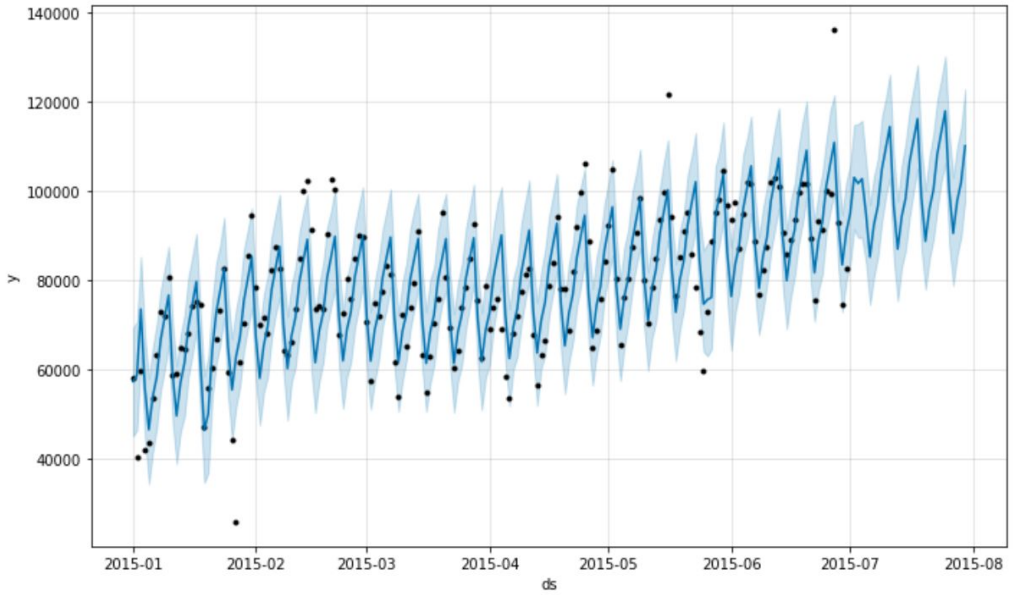

In [16]:
```
model.plot_components(forecast)
```

Out[16]:

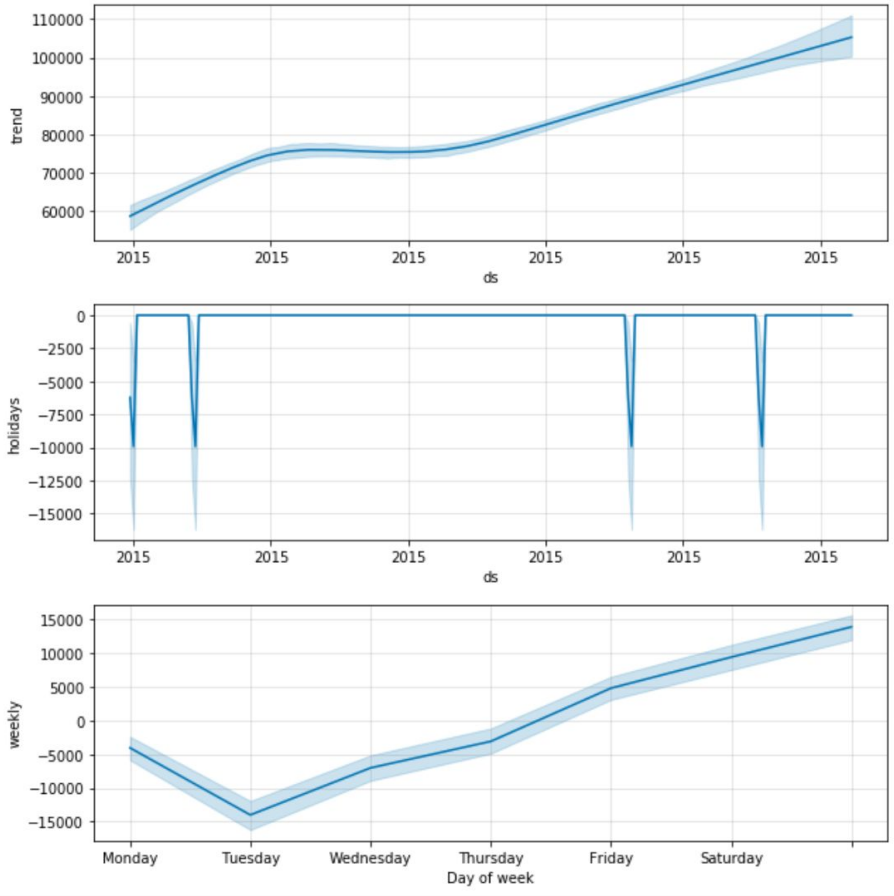

Sam Leach
github / twitter: samuelleach

https://github.com/samuelleach/python-notebooks